

Lab 3: PyTorch Operators and Intro to CNN

University of Washington

EE 596/AMATH 563

Spring 2021


Outline

- Python Operators
- Designing Training Procedures
- Introduction to Convolutional Neural Nets (CNNs)
- Exercise: MNIST Classification with CNN

PyTorchOperators and Layers


PyTorch Operators/Layers

- Activation Functions
- Normalization
- Dropout
- Loss Functions




PyTorch Operators/Layers

- Activation Functions
- Normalization
- Dropout
- Loss Functions




Activation Functions

- Non-linear functions performed by neurons
- ReLU - Rectified Linear Unit (nn.ReLU)
 - $y \geq 0$
- Tanh (nn.tanh)
 - $-1 < y < 1$
 - nn.Tanh
- Sigmoid (nn.Sigmoid)
 - $0 < y < 1$



Activation Functions


- Leaky ReLU
 - Similar to ReLU, but has non-zero values for negative x
 - Takes argument *negative_slope*, which determines the slope for $x < 0$.
- For full list of activation functions, see: <https://pytorch.org/docs/stable/nn.html>



Leaky ReLU with negative slope = 0.1

Python Operators/Layers

- Activation Functions
- **Normalization**
- Dropout
- Loss Functions




Input Normalization: Batch Normalization

- Normalizes input into each layer for each training mini-batch
- Addresses issue of shifting input distributions over training
- Inputs:
 - num_features: Number of features in the input vector
 - eps: numerical stability parameter

Example:

```
b = torch.nn.BatchNorm1D(100)
input = torch.rand(50, 100)
output = b(input)
```




Input Normalization

Other normalization procedures include:

- Layer Norm: Transposes Batch Norm. Normalizes over all summed inputs to a layer
 - <https://arxiv.org/abs/1607.06450>
- Group Norm: Normalizes by grouped channels instead of batches
 - <https://arxiv.org/abs/1803.08494>


Python Operators/Layers

- Activation Functions
- Normalization
- **Dropout**
- Loss Functions




Dropout

- Randomly zeroes some elements of input tensor with probability p
- Effective technique for regularization
- Outputs scaled by $1/1-p$
- Treated as identity during evaluation




Python Operators/Layers

- Activation Functions
- Normalization
- Dropout
- Loss Functions



Loss Functions

- Loss function parameters:
 - Reduction: how the output will be reduced in dimension:
 - None: Gives entire Loss Tensor with no reduction over batches
 - Sum: Takes sum of the loss tensor across batches, returning a single number
 - Mean: Same as sum, but divides by the number of batches to get the mean



Loss Functions – Cross Entropy Loss



- Cross Entropy Parameters

- Weight

- 1D tensor assigning weights to each class, which is helpful if you have an unbalanced training set

- ignore_index

- Specifies a target value that is ignored and does not contribute to the input gradient






Designing Training Procedures

Optimizer Initialization

- Parameters:
 - Should be iterable containing parameters to optimize
 - E.g., `model.parameters()` or `[var1, var2]`
 - Parameters must be defined BEFORE the optimizer
- Optimizer Settings
 - Learning rate, weight decay, etc.




Optimizers


- Stochastic Gradient Descent (`torch.optim.SGD`)
 - `params`: Model parameters
 - `lr`: Learning rate (required)
 - `momentum`: momentum factor (default: 0)
 - `weight_decay`: (default: 0)

Example: `torch.optim.SGD(model.parameters(), lr = 0.001, momentum = 0.2, weight_decay = 0.1)`

SGD without momentum



SGD with momentum



Optimizers

- Adam (`torch.optim.Adam`)
 - params: Model parameters
 - lr: Learning rate (default: 0.001)
 - betas: coefficients (tuple) used for computing running averages of gradient and its square (default: (0.9, 0.999))
 - eps: term added to denominator to improve numerical stability (default: 1e-8)
 - weight_decay: (default: 0)

Example:

```
torch.optim.Adam(model.parameters(),  
lr = 0.01, betas = (0.95, 0.998),  
eps = 1e-7)
```


Other Common Optimizers

- AdaDelta (`torch.optim.Adadelta`)
 - Precursor to Adam which uses first-order estimates to adapt learning rate
- Adamax (`torch.optim.Adamax`)
 - Variant on Adam based on infinity norm
- RMSProp (`torch.optim.RMSprop`)
 - Take the square root of the gradient average before adding epsilon to normalization of LR

Introduction to Convolutional Neural Nets (CNNs)


Convolutional Layers

- Convolutional Layers use a moving frame (kernel) to process input
- Helpful for learning local features
- Parameters of the kernel are learned during training




Convolutional Layers: Arguments

- `in_channels`: number of channels in the input image (e.g., RGB)
- `out_channels`: Number of output channels
- `kernel size`: Tuple (or int) indicating the dimensions of the convolving kernel



Convolutional Layers: Arguments

- stride: step size between each convolution – how far the kernel moves in each direction between convolutions (default: 1)
- padding- zero-padding added to each side of the input (default: 0)



Pooling


- Operates like Convolutional Layers, but perform simple math operations
- Moving frame calculates one of:
 - Max
 - Mean
 - Power-average (power defined by argument `norm_type`)

Example:


```
max_pool = nn.MaxPool2D(kernel_size = 2,  
stride= 2)
```

```
avg_pool = nn.AvgPool2D(kernel_size = 2,  
stride= 2)
```

Max Pooling



Avg Pooling



Example: CNN Implementation

Following Tutorial here: <https://medium.com/swlh/pytorch-real-step-by-step-implementation-of-cnn-on-mnist-304b7140605a>

Data Preparation

- Use `train_test_split` to create a validation set from your training data (should be in array/tensor form) (20%)

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_cv, y_train, y_cv = train_test_split(features, labels,
3                                               test_size = 0.2,
4                                               random_state = 1212)
5
6 X_train = np.array(X_train).reshape(33600, 784) #(33600, 784)
7 X_cv = np.array(X_cv).reshape(8400, 784) #(8400, 784)
```

Data Preparation

- Reshape data so that it has correct dimensions for CNN.
- MNIST images are grayscale, so they have only one channel
- Images should be 28x28

```
1 #Formatting on training set
2 train_x = X_train.reshape(33600, 1, 28, 28)
3 train_x = torch.from_numpy(train_x).float()
4 # converting the target into torch format
5 y_train = torch.from_numpy(np.array(y_train))
6 # shape of training data
7 train_x.shape, y_train.shape
8
9 #Formatting on testing set
10 X_cv = X_cv.reshape(8400, 1, 28, 28)
11 X_cv = torch.from_numpy(np.array(X_cv)).float()
12 # converting the target into torch format
13 y_cv = torch.from_numpy(np.array(y_cv))
14 X_cv.shape, y_cv.shape
```

Data Preparation

- Set batch size and create DataLoader objects for easier training

[illegible]

Model Definition: Initialization

- We will use a standard architecture consisting of:
 - Two 2d convolutional layers w/ filter size (3x3). Each layer has 16 and 32 output channels, respectively
 - Two 2d MaxPool layers with filter size (2x2)
 - ReLU activations
 - An FC layer of 800 nodes

```
1 # Create CNN Model
2 class CNNModel(nn.Module):
3     def __init__(self):
4         super(CNNModel, self).__init__()
5
6         # Convolution 1
7         self.cnn1 = nn.Conv2d(in_channels=1, out_channels=16,
8                                kernel_size=3, stride=1, padding=0)
9         self.relu1 = nn.ReLU()
10
11        # Max pool 1
12        self.maxpool1 = nn.MaxPool2d(kernel_size=2)
13
14        # Convolution 2
15        self.cnn2 = nn.Conv2d(in_channels=16, out_channels=32,
16                               kernel_size=3, stride=1, padding=0)
17        self.relu2 = nn.ReLU()
18
19        # Max pool 2
20        self.maxpool2 = nn.MaxPool2d(kernel_size=2)
21
22        # Fully connected 1
23        self.fc1 = nn.Linear(32 * 5 * 5, 10)
```

Model Definition: Forward()

- As the input gets processed at each step, the dimension of the images changes
- The convolutional layers increase the number of channels used to represent the data
- Output is logits

```
25     def forward(self, x):
26         # Input x dimensions:  #nx1x28x28
27         # Set 1
28         out = self.cnn1(x)      #nx16x26x26
29         out = self.relu1(out)
30         out = self.maxpool1(out)#nx16x13x13
31
32         # Set 2
33         out = self.cnn2(out)    #nx32x11x11
34         out = self.relu2(out)
35         out = self.maxpool2(out)#nx32x5x5
36
37         #Flatten
38         out = out.view(out.size(0), -1) #nx800
39
40         #Dense
41         out = self.fc1(out)     #nx10
42
43         return out
```

Define Hyperparameters, Loss, Optimizers

- Define training iterations, learning rate
- Classification- use CrossEntropyLoss
- For optimizer, we use SGD for this example

```
1 #Definition of hyperparameters
2 n_iters = 2500
3 num_epochs = n_iters / (len(train_x) / batch_size)
4 num_epochs = int(num_epochs)
5
6 # Cross Entropy Loss
7 error = nn.CrossEntropyLoss()
8
9 # SGD Optimizer
10 model = CNNModel()
11 learning_rate = 0.001
12 optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```


Model Training

- Define the quantities you want to track before the training loop
- `tensor.view()` allows you to reshape your input so that it is in the form your network needs

```
1 # CNN model training
2 count = 0
3 loss_list = []
4 iteration_list = []
5 accuracy_list = []
6 for epoch in range(num_epochs):
7     for i, (images, labels) in enumerate(train_loader):
8
9         train = Variable(images.view(100,1,28,28))
10        labels = Variable(labels)
11        # Clear gradients
12        optimizer.zero_grad()
13        # Forward propagation
14        outputs = model(train)
15        # Calculate softmax and cross entropy loss
16        loss = error(outputs, labels)
17        # Calculating gradients
18        loss.backward()
19        # Update parameters
20        optimizer.step()
21
22        count += 1
```

Model Training: Tracking progress

- Within training loop, track your accuracy on the validation/test set.
- Test and print at pre-defined intervals
- Track relevant information in lists defined above (loss, accuracy, iteration)

```
if count % 50 == 0:
    # Calculate Accuracy
    correct = 0
    total = 0
    # Iterate through test dataset
    for images, labels in test_loader:

        test = Variable(images.view(100,1,28,28))
        # Forward propagation
        outputs = model(test)
        # Get predictions from the maximum value
        predicted = torch.max(outputs.data, 1)[1]



        # Total number of labels
        total += len(labels)
        correct += (predicted == labels).sum()

    accuracy = 100 * correct / float(total)

    # store loss and iteration
    loss_list.append(loss.data)
    iteration_list.append(count)
    accuracy_list.append(accuracy)
if count % 500 == 0:
    # Print Loss
    print('Iteration: {} Loss: {} Accuracy: {} %'.format(count, loss.data, accuracy))
```


Plot your tracked quantities

```
1 # visualization loss
2 plt.plot(iteration_list,loss_list)
3 plt.xlabel("Number of iteration")
4 plt.ylabel("Loss")
5 plt.title("CNN: Loss vs Number of iteration")
6 plt.show()
7
8 # visualization accuracy
9 plt.plot(iteration_list,accuracy_list,color = "red")
10 plt.xlabel("Number of iteration")
11 plt.ylabel("Accuracy")
12 plt.title("CNN: Accuracy vs Number of iteration")
13 plt.show()
```



Exercise: MNIST Classification using LeNet-5

LeNet-5 Model




- Well-known network with seven layers, three of which are convolutional.

Image Source:

<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>


LeNet-5 Model Layers

- Layer 1: Convolutional Layer with 6 filters (output channels), kernel size of 5x5, and padding of 2
- Layer 2: Average pooling (2x2 kernel)
- Layer 3: Convolutional layer. 16 filters, 5x5 kernel size, no padding



LeNet Model Layers

- Layer 4: Average pooling (2x2)
- Layer 5: 120 filters of size 5x5. Output is 1x1x120
- Layer 6: Fully connected layer. Input dimensions: 120, Output dimensions: 84
- Layer 7: Fully-connected layer. Input dimensions: 84, Output dimensions: 120



Assignment Details

- Implement LeNet in PyTorch
- tanh activation
- Use Adam Optimizer
- Should be able to achieve greater than 95% accuracy

```
1 class LeNet5(nn.Module):
2
3     def __init__(self, n_classes):
4         super(LeNet5, self).__init__()
5
6         #define LeNet5
7
8     def forward(self, x):
9         #Define forward pass
10        return logits    #can also return probabilities
11                          #by performing softmax
```